

Redes de Computadores

LEIC-A

3 – Transport Layer (part 2)

Prof. Paulo Lobato Correia
IST, DEEC – Área Científica de Telecomunicações

1

Objectives

Application
Transport
Network
Link
Physical

- Understand the principles behind transport layer services:
 - Multiplexing/demultiplexing;
 - Reliable data transfer;
 - Flow control;
 - Congestion control.
- Transport layer protocols in the Internet:
 - UDP: connectionless transport;
 - TCP: connection-oriented transport with congestion control;
 - QUIC: Quick UDP Internet Connections.

4

Internet Transport Layer Protocols

Reliable, in-order delivery (TCP):


- Connection setup;
- Flow control;
- Congestion control.

Unreliable, unordered delivery (UDP):

- Simple extension of “best-effort” IP.

Services not available:

- Delay guarantees;
- Bandwidth guarantees.



RC – Prof. Paulo Lobato Correia 5

5

UDP: User Datagram Protocol

Often used for streaming multimedia applications:

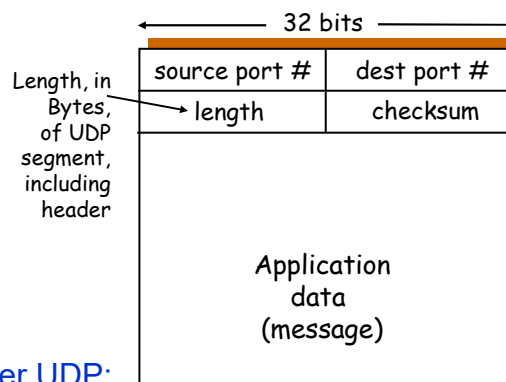
- Loss tolerant;
- Rate sensitive.

Other UDP uses

- DNS;
- SNMP;
- QUIC.

To have reliable transfer over UDP:

- Reliability added at application layer;
- Application-specific error recovery!




UDP segment format

RC – Prof. Paulo Lobato Correia 6

6

Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- **Connection-oriented transport: TCP**
 - **Connection management;**
 - **Segment structure;**
 - **Reliable data transfer;**
 - **Flow control;**
- Principles of congestion control
- TCP congestion control
- QUIC



RC – Prof. Paulo Lobato Correia 7

7

TCP: Overview

Some RFCs: 675, **793**, **1122**, 1323, 1379, 1948, 2018, **5681**, 6247, **6298**, 6824, **7323**, 7424, ...

- RFC 793: "Transmission Control Protocol", STD 7 (Sept 1981) - fundamental TCP specification document
- RFC 1122: "Requirements for Internet Hosts - Communication Layers" (Oct 1989) - updates and clarifies RFC 793, fixes specification bugs and oversights. **Mandates that a congestion control mechanism must be implemented.**
- RFC 5681: "TCP Congestion Control" (Aug 2009) - defines **congestion avoidance and control mechanism for TCP**. RFCs 2001 and 2581 are conceptual precursors of RFC 5681.
- RFC 6298: "Computing TCP's Retransmission Timer" (June 2011)
- RFC 6691: "TCP Options and Maximum Segment Size (MSS)" (July 2012)
- RFC 7323: "TCP Extensions for High Performance" (Sept 2014)
- RFC 7414 - A Roadmap for Transmission Control Protocol (TCP) Specification Documents

RC – Prof. Paulo Lobato Correia 8


8

TCP: Overview

- **Point-to-point:**
 - One sender, one receiver;
- **Reliable, in-order *byte stream*:**
 - No “message boundaries”;
- **Sliding window:**
 - TCP congestion and flow control set window size;
- **Send & receive buffers;**
- **Full duplex data:**
 - Bi-directional data flow in same connection;
 - **MSS**: maximum segment size;
- **Connection-oriented:**
 - Handshaking (exchange of control messages)
initializes sender and receiver state before data exchange;
- **Flow control:**
 - Sender will not overwhelm receiver.



TCP segment structure



TCP Connection Management

Recall:

- TCP sender and receiver establish "connection" before exchanging data segments;
- Initialize TCP variables: seq. numbers, buffers, flow control information (e.g. **RcvWindow**);
- *Client*: initiates connection;
- *Server*: contacted by client.

Three-way handshake:


Step 1: Client sends TCP **SYN** segment to server:

- Specifies initial client seq. number (x);
- No data included.

Step 2: Server receives SYN and replies with **SYN ACK** segment:

- Server allocates buffers;
- Specifies server initial seq. number (y).

Step 3: Client receives SYN ACK, replies with **ACK** segment, which **may already contain data**.



RC – Prof. Paulo Lobato Correia 11

11

TCP Closing a Connection

Client and server each close their side of connection.

Example: client closes socket (*server could start*) :


Step 1: Client end system sends TCP **FIN** control segment to server;

Step 2: Server receives FIN, replies with **ACK**. Closes connection, sends **FIN**.

Step 3: Client receives FIN, replies with **ACK**.

- Enters "**timed wait**" - will respond with ACK to received FINs.

Step 4: Server, receives ACK. Connection closed.




RC – Prof. Paulo Lobato Correia 12

12

Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- **Connection-oriented transport: TCP**
 - Connection management;
 - **Segment structure;**
 - Reliable data transfer;
 - Flow control;
- Principles of congestion control
- TCP congestion control
- QUIC



RC – Prof. Paulo Lobato Correia 14

14

TCP Sequence Numbers, ACKs

Sequence numbers:

- Byte stream – use **“number” of first byte** in segment’s data

Acknowledgements:


- Seq # of next byte expected from other side
- Cumulative ACK

outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
A	rwnd
checksum	urg pointer

outgoing segment from sender

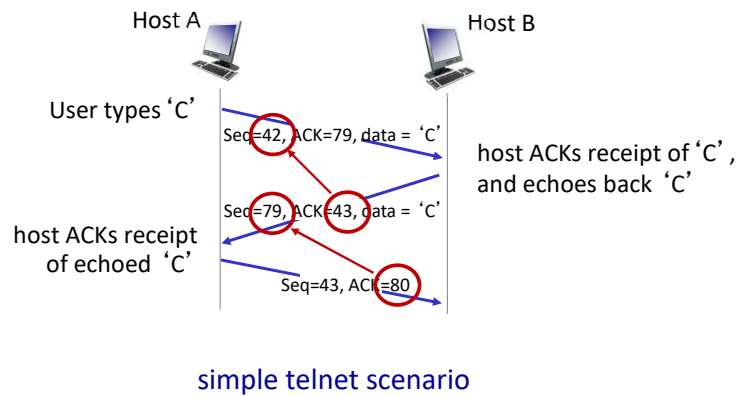
source port #	dest port #
sequence number	
acknowledgement number	
checksum	urg pointer



RC – Prof. Paulo Lobato Correia 16

16

TCP Sequence Numbers, ACKs



TCP Round Trip Time and Timeout

How to set TCP timeout value?

- Longer than RTT:
 - But RTT varies...
- Too short: premature timeout – unnecessary retransmissions.
- Too long: slow reaction to segment loss.


How to estimate RTT?

- **SampleRTT**: measured time between segment transmission and ACK receipt.
 - Ignore retransmissions;
- **SampleRTT** varies; Smoother RTT estimation (**EstimatedRTT**):
 - Average several recent measurements, not just current **SampleRTT**.

TCP Round Trip Time, Timeout

$$\text{EstimateRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponential weighted moving average (EWMA)
- Influence of past sample decreases exponentially fast
- Typical value: $\alpha = 0.125$



RC – Prof. Paulo Lobato Correia 19

19

TCP Round Trip Time, Timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
 - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑
estimated RTT

↑
“safety margin”

- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)


* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

RC – Prof. Paulo Lobato Correia 20

20

Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- **Connection-oriented transport: TCP**
 - Connection management;
 - Segment structure;
 - **Reliable data transfer**;
 - Flow control;
- Principles of congestion control
- TCP congestion control
- QUIC



RC – Prof. Paulo Lobato Correia 21

21

TCP Reliable Data Transfer

- TCP creates a reliable data transfer (rdt) service on top of the unreliable service offered by IP;
- TCP uses:
 - **Sliding window**;
 - **Cumulative ACKs**;
 - A single **retransmission timer**;
- Retransmissions are triggered by:
 - Timeout events;
 - Duplicate ACKs.

Initially, let's consider a simplified TCP sender:

- Ignore duplicate ACKs
- Ignore flow control and congestion control.

RC – Prof. Paulo Lobato Correia 22

22

TCP Sender Events

Data received from the application layer:

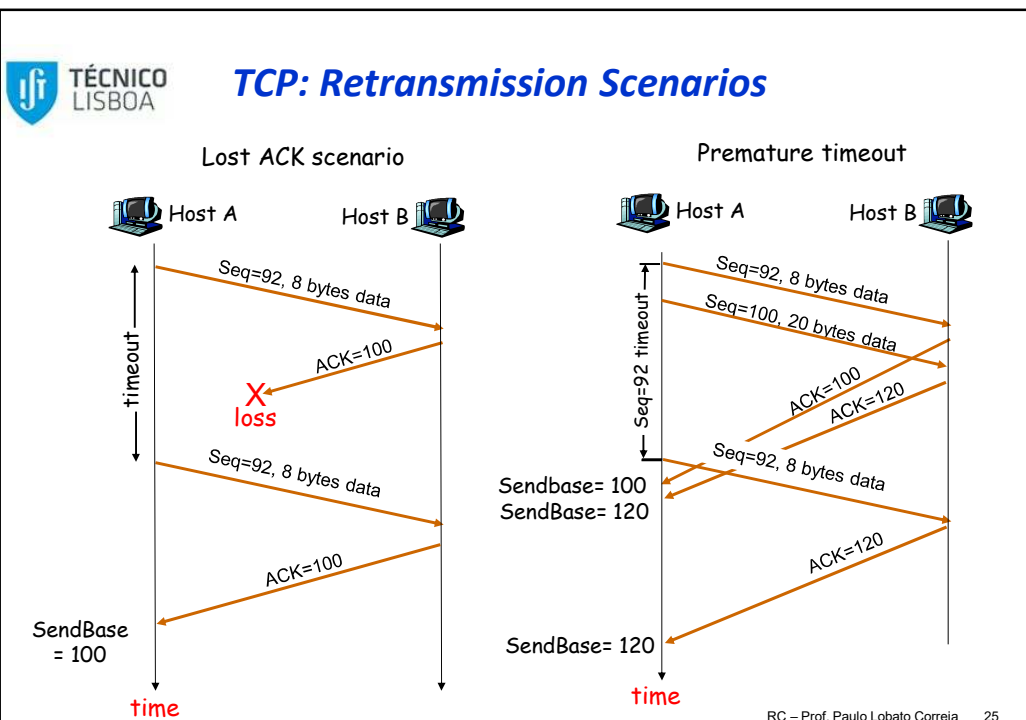
- Create segment;
- Sequence number is the **number** of the **first data byte** in the segment;
- Start timer (if not yet running – timer is for the oldest unacked segment);
- Expiration interval: **TimeoutInterval**

Timeout:

- Retransmit the segment that caused the timeout;
- Restart the timer

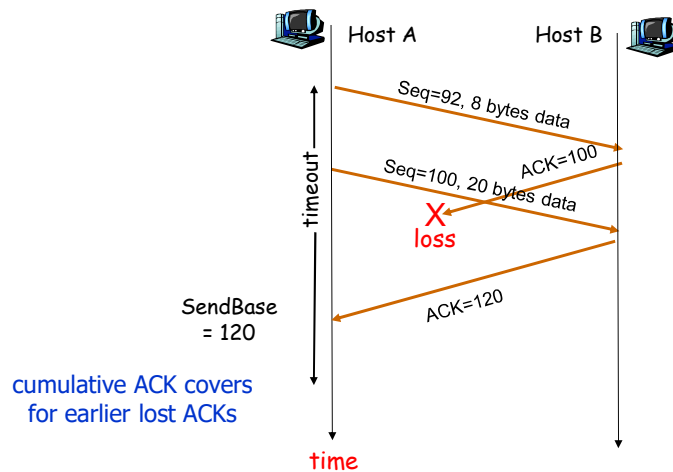
ACK received:

- If previously unacked, segments are acknowledged:
 - Update what is known to be ACKed;
 - Start timer if there are other outstanding segments.



TCP: Retransmission Scenarios

Cumulative ACK scenario



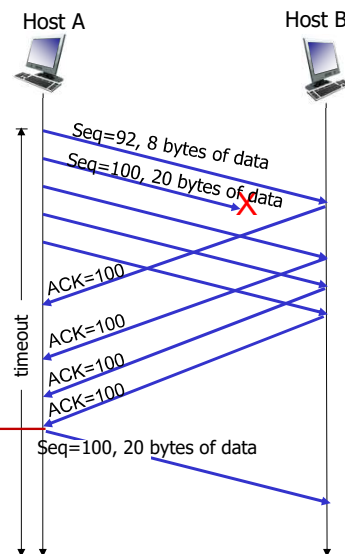
26

TCP Fast Retransmit

- Time-out period is relatively long:
- **Detect lost segments via duplicate ACKs.**
 - Sender often sends many data segments;
 - If a segment is lost → many duplicate ACKs.
- If sender receives 3 additional ACKs, assumes the segment (after the ACKed data) was lost:
 - **Fast retransmit:**
resends segment before timer expiration.



Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. Retransmit!




27

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK . Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK , ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK , indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediately send ACK , provided that segment starts at lower end of gap

Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- **Connection-oriented transport: TCP**
 - Connection management;
 - Segment structure;
 - Reliable data transfer;
 - **Flow control**;
- Principles of congestion control
- TCP congestion control
- QUIC




TCP Flow Control

Goal: ensure that sender won't overflow the receiver's buffer by transmitting too fast.

TCP has a receive buffer:

Application process may be slow at reading from buffer


Speed-matching: adjust send rate to the receiving application processing rate



TCP Flow Control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - many operating systems autoadjust **RcvBuffer**
- Sender limits amount of unACKed (“in-flight”) data to received **rwnd**
 - Guarantees receive buffer will not overflow

flow control: # bytes receiver willing to accept




TCP segment format

Test applet at:

https://media.pearsoncmg.com/aw/ecs_kurose_compnetwork_7/cw/content/interactiveanimations/flow-control/index.html

Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- Connection-oriented transport: TCP
 - Connection management;
 - Segment structure;
 - Reliable data transfer;
 - Flow control;
- Principles of congestion control
- TCP congestion control



RC – Prof. Paulo Lobato Correia 34

34

Principles of Congestion Control

Congestion:

- Informally:
 - “Too many sources sending **too much data too fast** for the *network* to handle”;
- Different from flow control!
- Manifestations/consequences:
 - **Lost packets** (*buffer overflow at routers*);
 - **Long delays** (*queueing in router buffers*).
- A top-10 problem!



congestion control:
too many senders,
sending too fast


RC – Prof. Paulo Lobato Correia 35

35



Causes/Costs of Congestion

- four senders
- multi-hop paths
- timeout/retransmit

As red λ_{in}' increases, all arriving blue pkts at upper queue are dropped, blue throughput $\rightarrow 0$



Causes/Costs of Congestion




Another “cost” of congestion:

- when packet dropped, any upstream transmission capacity and buffering used for that packet was wasted!

Causes/costs of Congestion: Insights

- Throughput can never exceed capacity
- Delay increases as capacity approached
- Loss/retransmission decreases effective throughput
- Un-needed duplicates further decreases effective throughput
- Upstream transmission capacity / buffering wasted for packets lost downstream



RC – Prof. Paulo Lobato Correia 47

47

Approaches Towards Congestion Control

Two approaches towards congestion control:

Network-assisted congestion control:

- Routers provide feedback to end systems:
 - Bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM);
 - Router may inform sender explicitly of supported rate.

End-end congestion control:

- No explicit feedback from network;
- Congestion inferred from end-system observed loss and delay;
 - It's the approach taken by TCP.


RC – Prof. Paulo Lobato Correia 48

48

Network-assisted congestion control:


- Routers provide *direct* feedback to sending/receiving hosts with flows passing through congested router
- May indicate congestion level or explicitly set sending rate

□ TCP ECN, ATM, DECbit protocols




End-end congestion control:

- No explicit feedback from network
- Congestion *inferred* from observed loss, delay
- Approach taken by TCP



Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- Connection-oriented transport: TCP
 - Connection management;
 - Segment structure;
 - Reliable data transfer;
 - Flow control;
- Principles of congestion control
- TCP congestion control
- QUIC



RC – Prof. Paulo Lobato Correia 53

53

TCP Congestion Control: AIMD


- **Approach:** senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event

Additive Increase

Increase sending rate by 1 maximum segment size every RTT until loss detected

Multiplicative Decrease

Cut sending rate in half at each loss event



AIMD sawtooth behavior: *probing* for bandwidth

RC – Prof. Paulo Lobato Correia 54

54

TCP AIMD: more


Multiplicative decrease detail – sending rate is:

- Cut in half on loss detected by triple duplicate ACK (TCP Reno)
- Cut to 1 MSS (maximum segment size) when loss detected by timeout (TCP Tahoe)

Why AIMD?

- AIMD – a distributed, asynchronous algorithm – has been shown to:
 - Optimize congested flow rates network wide!
 - Have desirable stability properties

TCP Congestion Control: Details



TCP sending behavior:

- *Roughly:* send `cwnd` bytes, wait RTT for ACKS, then send more bytes

$$\text{TCP rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- TCP sender limits transmission: $\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$
- `cwnd` is dynamically adjusted in response to observed network congestion (implementing TCP congestion control)

typically $\text{cwnd} < \text{rwnd}$

TCP Congestion Control: Details

How does sender perceive congestion?

- Loss events:
 - Timeout (TCP Tahoe);
 - 3 duplicate ACKs (TCP Reno).
- TCP sender reduces rate (**cwnd**) after loss event.


Three mechanisms:

- AIMD (additive increase, multiplicative decrease);
- Slow start;
- Conservative after timeout events.

TCP Slow Start


- When connection begins, rate increases exponentially until first loss event:
 - Initially **cwnd** = 1 MSS
 - Double **cwnd** every RTT;
 - This is done by incrementing **CongWin** of MSS for every ACK received:

$$cwnd = cwnd + MSS$$
- Slow start: The **initial rate is slow** but it **grows exponentially fast**.



TCP Slow Start

- When connection begins, **cwnd** = 1 MSS
 - Example: MSS = 500 bytes & RTT = 200 msec
 - Initial rate (\sim MSS/RTT) = 20 kbps
 - Too slow ...
- Available bandwidth may be \gg MSS/RTT
 - It is desirable to quickly ramp up to a respectable rate.
- **Slow Start:**
 - When connection begins,
rate increases **exponentially fast**
until first loss event or cwnd = Thr.



Loss Events

- After 3 duplicate ACKs (\rightarrow **Congestion Avoidance**):
 - **cwnd** is cut in half;
 - Window then grows linearly.
- But after *timeout* (\rightarrow **Slow Start**):
 - **cwnd** is set to 1 MSS;
 - The window then grows exponentially up to a threshold, then grows linearly.

Philosophy:

- 3 duplicate ACKs indicate that the network is still capable of delivering some segments;
- A timeout indicates a “more alarming” congestion scenario.


Loss Events

Q: When should the exponential increase switch to linear?

A: When **cwnd** gets to 1/2 of its value before timeout.

Implementation:

- Variable **Threshold**;
- At loss event, **Threshold** is set to 1/2 of **cwnd** just before loss event.




RC – Prof. Paulo Lobato Correia 61

61

Summary: TCP Congestion Control

- When **cwnd** is below **Threshold**, sender is in **slow-start** phase: the window grows exponentially.
- When **cwnd** is above **Threshold**, sender is in **congestion-avoidance** phase: window grows linearly.
- When a **triple duplicate ACK** occurs: **Threshold** set to **cwnd/2** and **cwnd** set to **Threshold**.
- When **timeout** occurs: **Threshold** set to **cwnd/2** and **cwnd** is set to 1 MSS.



RC – Prof. Paulo Lobato Correia 62

62

Fast Recovery state

When detecting **three duplicated ACKs** and before moving to congestion avoidance state, *try to speed up recovery* by fast resend of unacknowledged data:

- **cwnd** is increased by 1 MSS for every duplicate ACK received for the missing segment that caused TCP to enter the **fast recovery state**;

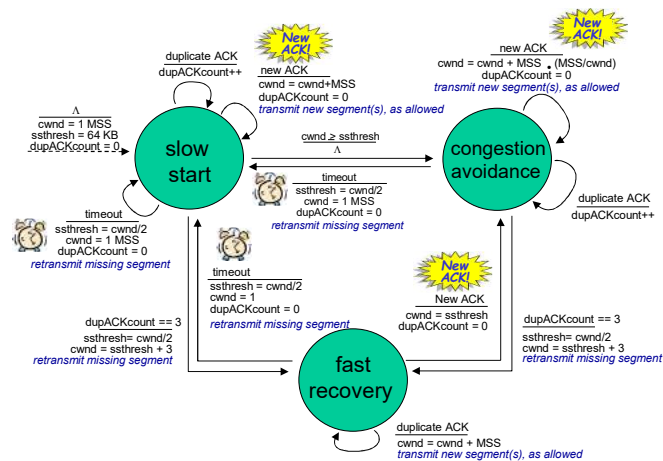
When receiving ACK for the missing segment:

- Move to **congestion avoidance** state (after deflating **cwnd**).

Fast recovery is a recommended, but not required, component of TCP [RFC 5681].

State	Event	TCP Sender Action	Comment
Slow Start (SS)	ACK receipt for previously unacked data	$cwnd = cwnd + MSS$, If ($cwnd > Threshold$) state = " Congestion Avoidance "	Double cwnd every RTT
Congestion Avoidance (CA)	ACK receipt for previously unacked data	$cwnd = cwnd + MSS * (MSS / cwnd)$	Additive increase: cwnd increases 1 MSS every RTT
SS or CA	Loss event detected by triple duplicate ACK	$Threshold = cwnd / 2$, $cwnd = Threshold + 3 MSS$, State = " Fast Recovery "	Enter fast recovery , implementing multiplicative decrease . cwnd will not drop below 1 MSS.
Fast Recovery (FR)	ACK receipt for the previously unacked data	$cwnd = Threshold$, $dupACKcount = 0$ State = " Congestion Avoidance "	Exit fast recovery
FR	duplicate ACK	$cwnd = cwnd + MSS$	FR – increase cwnd
SS, CA or FR	Timeout	$Threshold = cwnd / 2$, $cwnd = 1 MSS$, State = " Slow Start "	Enter Slow Start
SS or CA	Duplicate ACK	$dupACKcount ++$	cwnd and Threshold not changed

TCP Sender Congestion Control




RC – Prof. Paulo Lobato Correia 67

67

TCP CUBIC

- Is there a better way than AIMD to “probe” for usable bandwidth?
- Insight/intuition:
 - W_{max} : sending rate at which congestion loss was detected
 - congestion state of bottleneck link probably (?) hasn’t changed much
 - after cutting rate/window in half on loss:
 - start towards W_{max} **faster**, but then approach W_{max} **slowly**




TCP CUBIC - default in Linux, most popular TCP for popular Web servers

RC – Prof. Paulo Lobato Correia 68

68

TCP and Congested “bottleneck link”

- TCP (classic, CUBIC) increase TCP’s sending rate until packet loss occurs at some router’s output: the **bottleneck link**

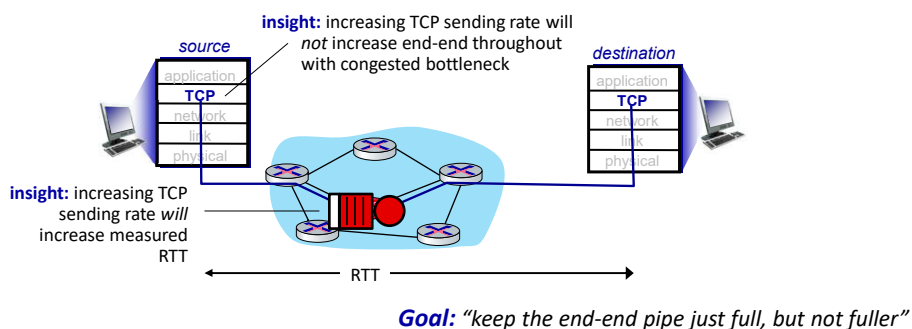


RC – Prof. Paulo Lobato Correia 70

70

TCP and Congested “bottleneck link”

- TCP (classic, CUBIC) increase TCP’s sending rate until packet loss occurs at some router’s output: the **bottleneck link**
- Understanding congestion: useful to focus on congested bottleneck link



RC – Prof. Paulo Lobato Correia 71

71

Delay-based TCP Congestion Control

Keeping sender-to-receiver pipe “just full enough, but no fuller”: keep bottleneck link busy transmitting, but avoid high delays/buffering



Delay-based approach:

- RTT_{min} - minimum observed RTT (uncongested path)
- uncongested throughput with congestion window cwnd is $\text{cwnd}/\text{RTT}_{\text{min}}$

if measured throughput “very close” to uncongested throughput
increase cwnd linearly /* since path **not congested** */

else if measured throughput “far below” uncongested throughput
decrease cwnd linearly /* since path is **congested** */


Delay-based TCP Congestion Control

- Congestion control without inducing/forcing loss
- Maximize throughput (“keeping the just pipe full...”) while keeping delay low (“...but not fuller”)
- A number of deployed TCPs take a delay-based approach:
 - BBR deployed on Google’s (internal) backbone network


Explicit congestion notification (ECN)

TCP deployments often implement *network-assisted* congestion control:

- two bits in IP header (ToS field) marked *by network router* to indicate congestion
 - policy* to determine marking chosen by network operator
- congestion indication carried to destination
- destination sets ECE bit on ACK segment to notify sender of congestion
- involves both IP (IP header ECN bit marking) and TCP (TCP header C,E bit marking)




TCP segment structure



TCP Fairness

Fairness goal:


For k TCP sessions sharing the same bottleneck link, of bandwidth R , each should have average rate of R/k .



Is TCP Fair?

Two competing sessions (both have a large amount of data to send, in CA mode):

- Additive increase gives slope of 1, as throughput increases;
- Multiplicative decrease reduces throughput proportionally.



Fairness

Fairness and parallel TCP connections

- Application can open parallel connections between 2 hosts;
- Web browsers do this.
- Example: link of rate R supporting 9 connections;
 - New app asks for 1 TCP, gets rate $R/10$
 - Another new app asks for 10 TCPs, getting $R/2$!


Fairness and UDP

- Multimedia applications often do not use TCP:
 - Do not want rate restricted by congestion control.
- Instead use UDP:
 - Send audio/video at constant rate and tolerate packet loss.

There is no “Internet police” to check the usage of congestion control.


Outline

- Transport-layer services;
- Multiplexing and demultiplexing;
- Connectionless transport: UDP;
- Principles of reliable data transfer;
- Connection-oriented transport: TCP
 - Connection management;
 - Segment structure;
 - Reliable data transfer;
 - Flow control;
- Principles of congestion control
- TCP congestion control
- QUIC



QUIC: Quick UDP Internet Connections

- Application-layer protocol, on top of UDP
 - Increase performance of HTTP
 - Deployed on many Google servers, apps (e.g., Chrome, mobile YouTube app)




QUIC: Quick UDP Internet Connections


QUIC:

- **Error and congestion control:**
 “Readers familiar with TCP’s loss detection and congestion control will find algorithms here that parallel well-known TCP ones.”
[from QUIC specification]
- **Connection establishment:**
 Reliability, congestion control, authentication, encryption, state – all established in just one RTT
- Multiple application-level “streams” multiplexed over single QUIC connection
 - Separate reliable data transfer, security
 - Common congestion control

QUIC: Connection Establishment




TCP (reliability, congestion control state)
+
TLS (authentication, crypto state)
▪ 2 serial handshakes



QUIC: reliability, congestion control, authentication, crypto state
▪ 1 handshake

82

QUIC Streams: Parallelism, no HOL Blocking



(a) HTTP 1.1

83

Chapter 3: Summary

- Principles behind transport layer services:
 - Multiplexing and demultiplexing;
 - Reliable data transfer;
 - Flow control;
 - Congestion control.
- Instantiation and implementation in the Internet:
 - UDP;
 - TCP;
 - QUIC.