

# Integrated Operation of the Memory System

## Computer Organization

Friday, 07 October 2022

Many slides adapted from:  
Computer Organization and Design,  
Patterson & Hennessy  
5th Edition, © 2014, MK  
and from Prof. Mary Jane Irwin, PSU

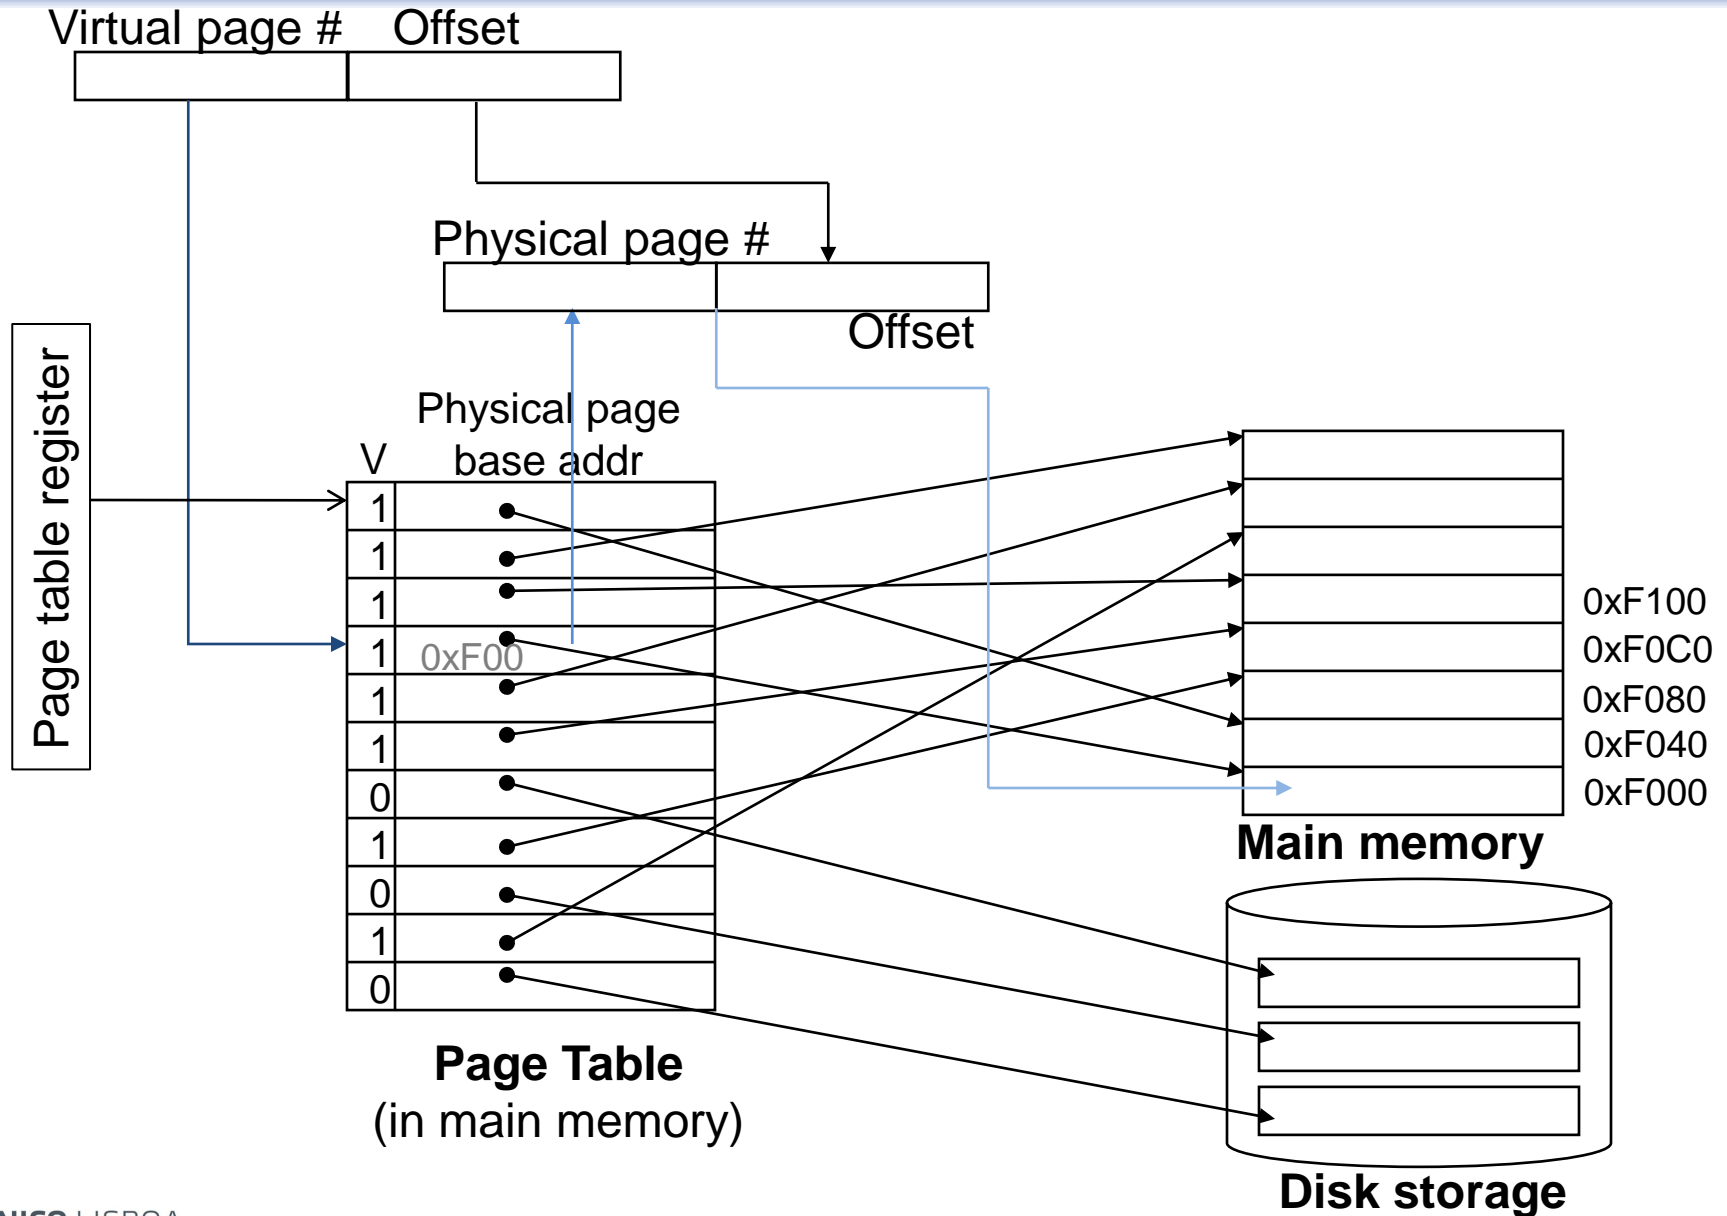


**TÉCNICO** LISBOA

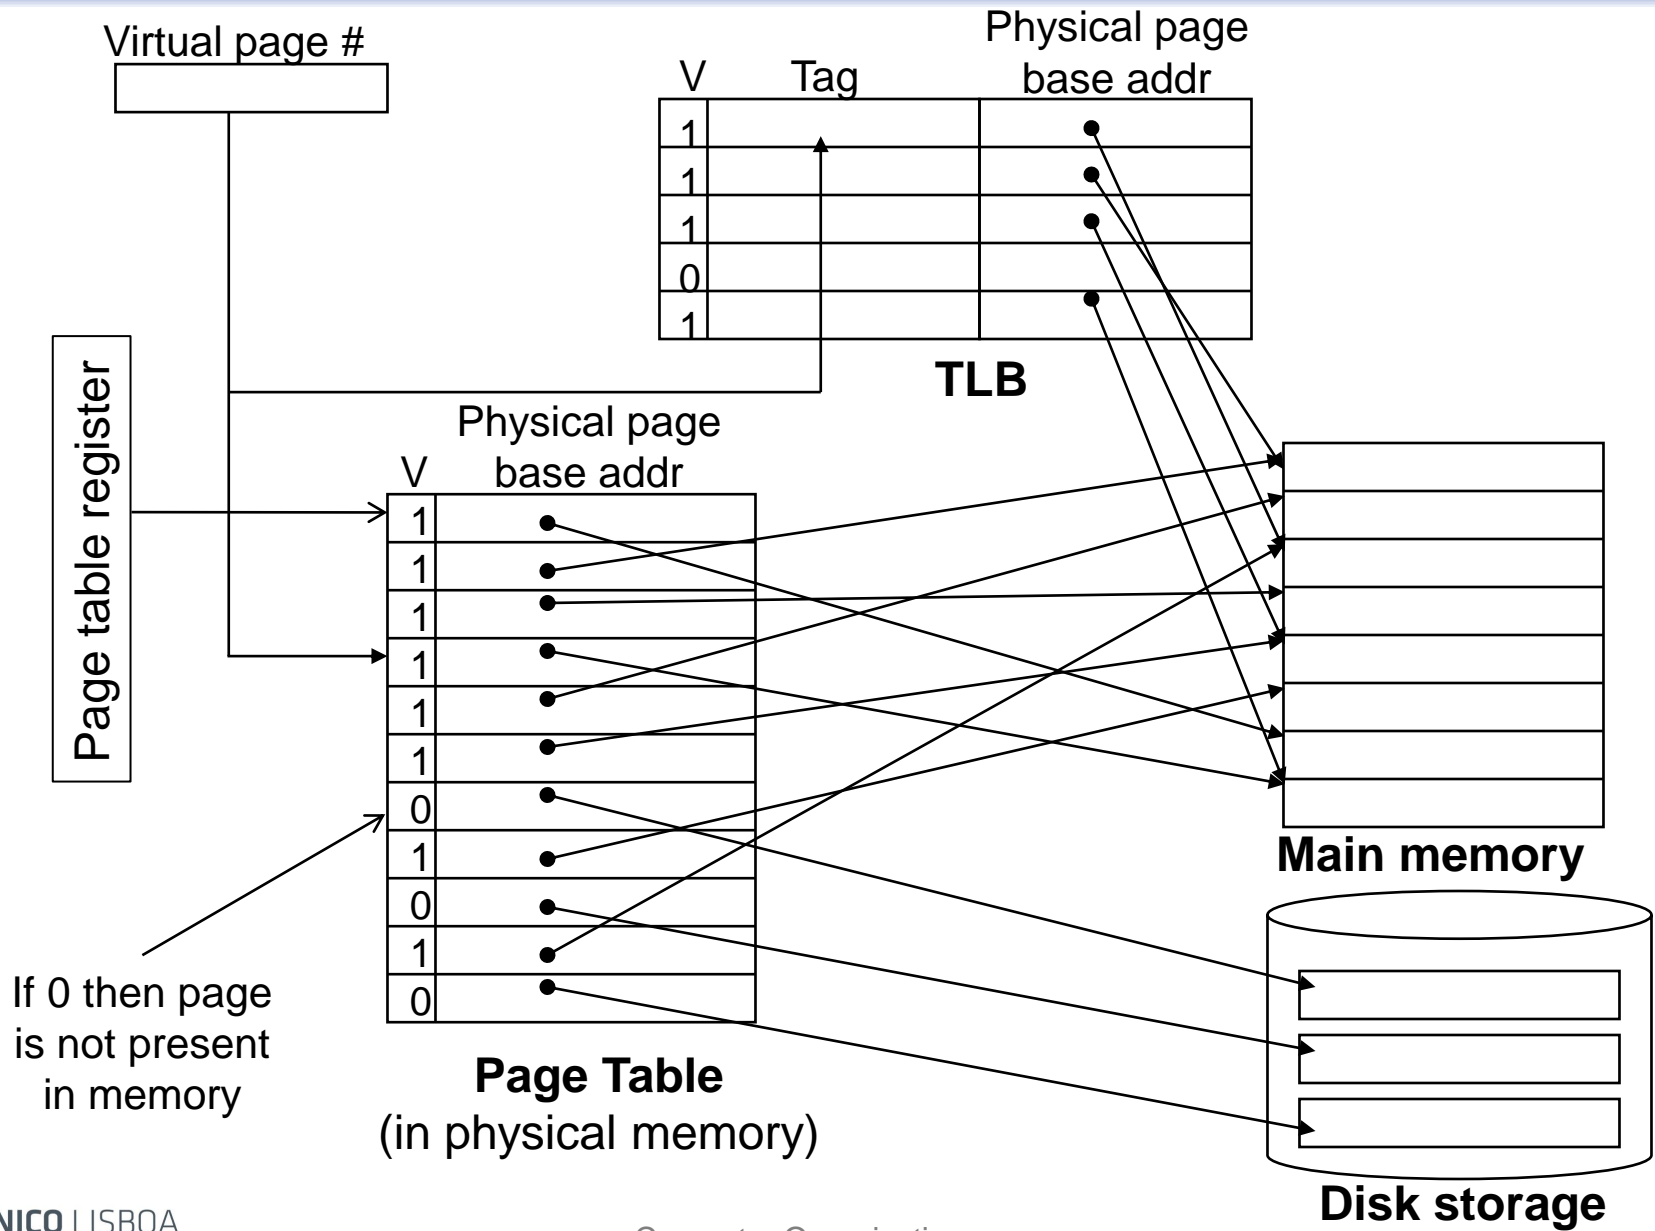
# Summary

- Previous Classes
  - Virtual Memory
- Today:
  - Translation Lookaside Buffer (TLB)
  - Integrated Operation of the Memory System

# Translation Using a Page Table



# Fast Translation Using a TLB



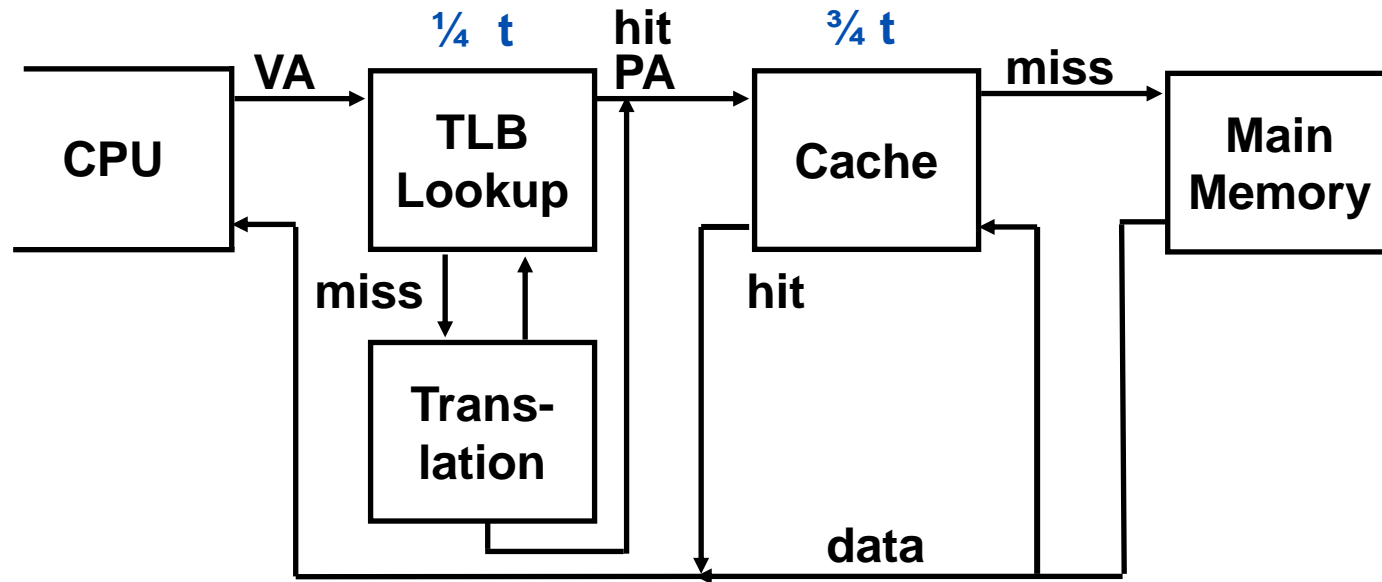
# Translation Lookaside Buffers (TLBs)

- Just like any other cache, the TLB can be organized as fully associative, set associative, or direct mapped

| V | Virtual Page # | Physical Page # | Dirty | Ref | Access |
|---|----------------|-----------------|-------|-----|--------|
|   |                |                 |       |     |        |

- TLB access time is typically smaller than cache access time (because TLBs are much smaller than caches)
  - Typical: 16–512 PTEs, 0.5–1 cycle for hit, 10–100 cycles for miss, 0.01%–1% miss rate

# A TLB in the Memory Hierarchy



- A TLB miss – is it a page fault or merely a TLB miss?
  - If the page is loaded into main memory, then the TLB miss can be handled (in hardware or software) by loading the translation information from the page table into the TLB
    - Takes 10's of cycles to find and load the translation info into the TLB
  - If the page is not in main memory, then it's a true page fault
    - Takes 1,000,000's of cycles to service a page fault
- TLB misses are much more frequent than true page faults

# TLB Miss Handler

- TLB miss indicates
  - Page present, but PTE not in TLB
  - Page not present
- Must recognize TLB miss before destination register overwritten
  - Raise exception
- Handler copies PTE from memory to TLB
  - Then restarts instruction
  - If page not present, page fault will occur

# Page Fault Handler


- Use faulting virtual address to find PTE
- Locate page on disk
- Choose page to replace
  - If dirty, write to disk first
- Read page into memory and update page table
- Make process runnable again
  - Restart from faulting instruction



# Check@home: TLB Event Combinations

| TLB  | Page Table | Cache        | Possible? Under what circumstances?  |
|------|------------|--------------|--|
| Hit  | Hit        | Hit          | Yes – what we want!  |
| Hit  | Hit        | Miss         | Yes – although the page table is not checked if the TLB hits               |
| Miss | Hit        | Hit          | Yes – TLB miss, PA in page table   |
| Miss | Hit        | Miss         | Yes – TLB miss, PA in page table, but data not in cache                    |
| Miss | Miss       | Miss         | Yes – page fault   |
| Hit  | Miss       | Miss/<br>Hit | Impossible – TLB translation not possible if page is not present in memory |
| Miss | Miss       | Hit          | Impossible – data not allowed in cache if page is not in memory            |

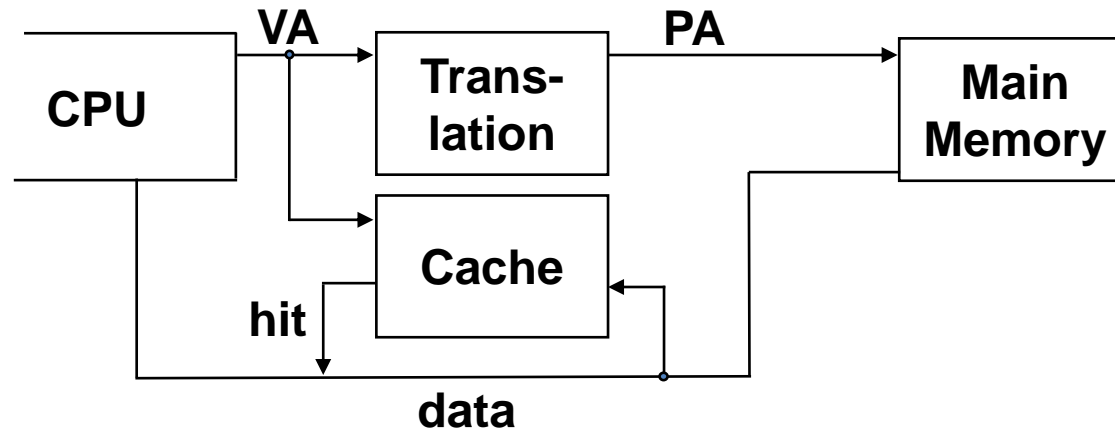
# TLB and Cache Interaction



- If cache tag uses physical address
  - Need to translate before cache lookup
- Alternative: use virtual address tag
  - Complications due to aliasing
    - Different virtual addresses for shared physical address

# Why Not a Virtually Addressed Cache?

- A virtually addressed cache would only require address translation on cache misses



But, two programs which are sharing data will have two different virtual addresses for the same physical address – aliasing – so have two copies of the shared data in the cache and two entries in the TLB which would lead to coherence issues

- Must update all cache entries with the same physical address or the memory becomes inconsistent

# Reducing Translation Time

- Interpretation of the virtual address by the TLB:

|                    |                |
|--------------------|----------------|
| Virtual page index | Virtual offset |
|--------------------|----------------|

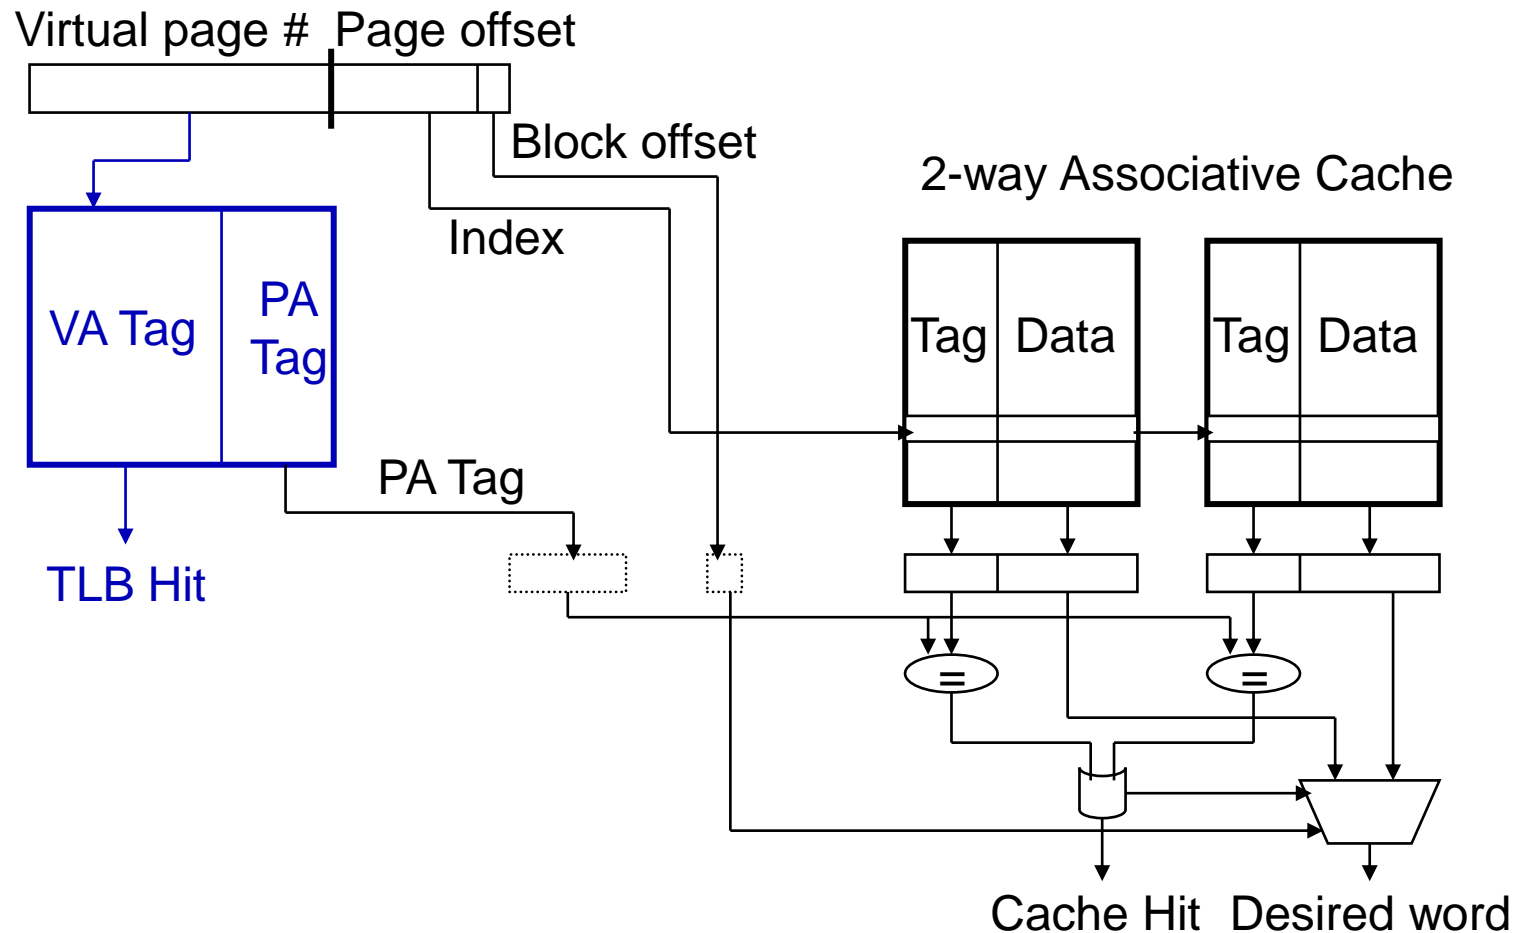
- Interpretation of the physical address by the cache:

|     |       |        |
|-----|-------|--------|
| Tag | Index | offset |
|-----|-------|--------|

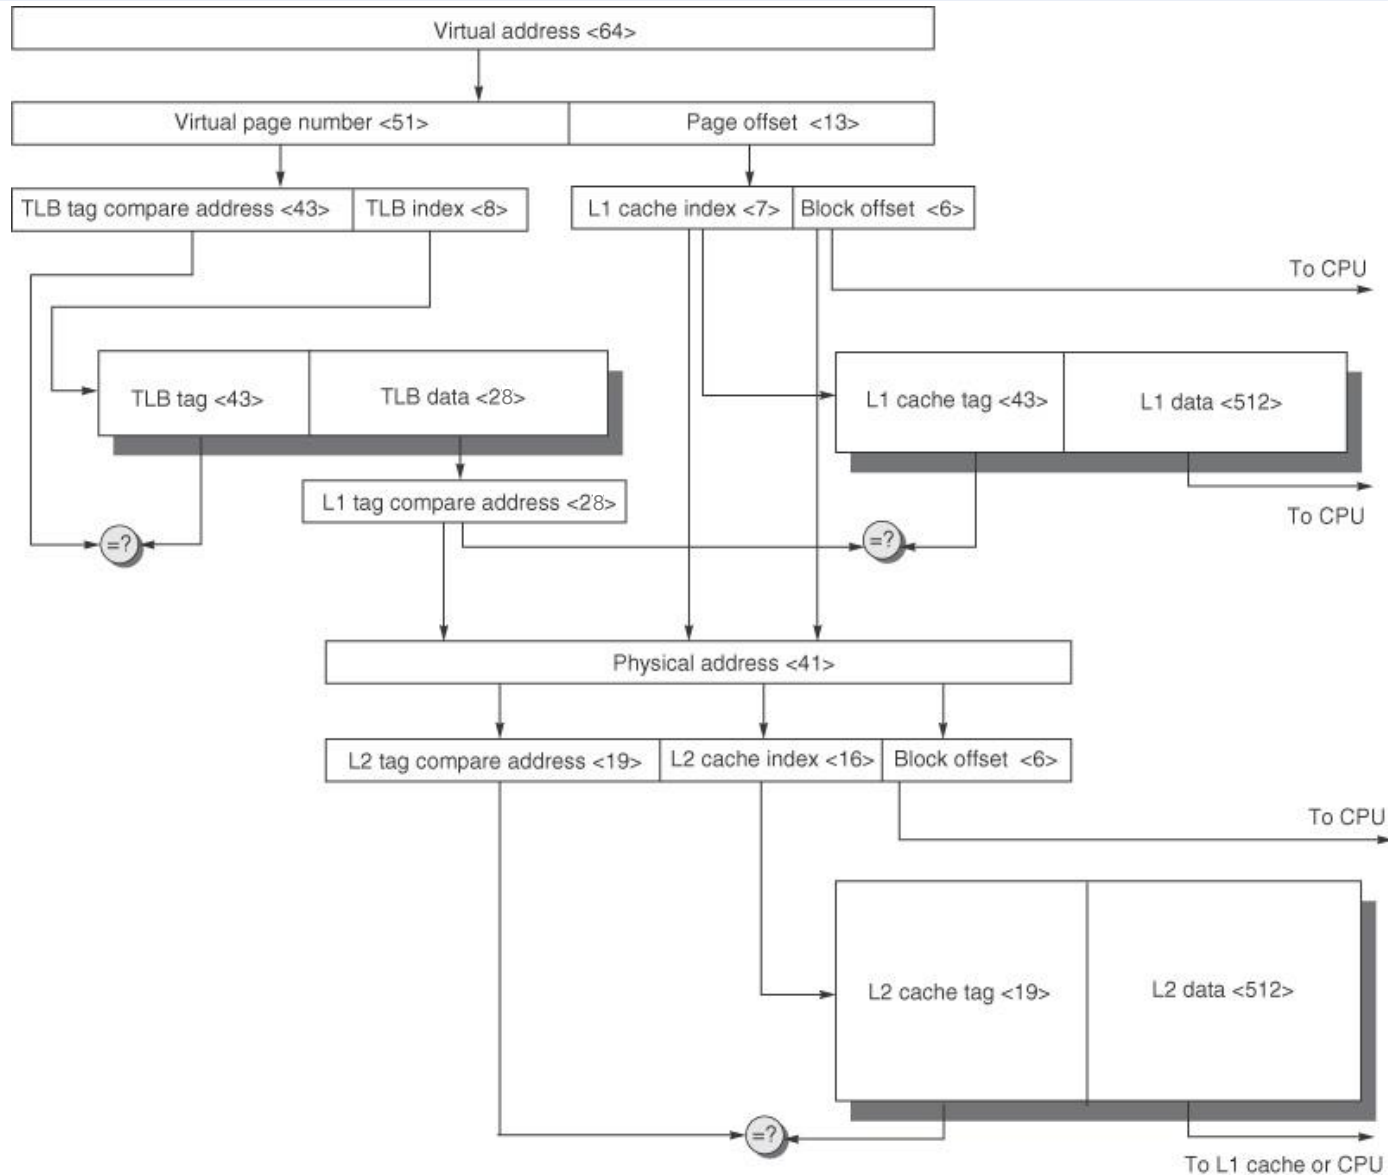
- How can we parallelize the access?
  - The virtual address offset field does not participate in the address translation process.
  - If Virtual offset  $\geq$  Index+offset, the cache index field is included within the virtual address offset.
- Cache may be read in parallel with the TLB test
- **L1 Cache:** “virtually indexed, physically tagged”

# Parallel Access Cache-TLB

- Can **overlap** the cache access with the TLB access
  - Works when the high order bits of the VA are used to access the TLB while the low order bits are used as index into cache



# Parallel Access Cache-TLB



# Parallel Access Cache-TLB

## Possible scenarios:

- Hit, both in TLB and in cache:
  - Access time is similar to the cache access
- Hit in TLB, Miss in cache:
  - Access time is similar to a primary memory access, with a cache miss
- Miss in TLB:
  - It is necessary to wait for the translation (by hierarchy or inverted table)
  - There is some gain in the cache access (but it is not significant...)

# The Hardware/Software Boundary

- What parts of the virtual to physical address translation is done by or assisted by the hardware?
  - Translation Lookaside Buffer (TLB) that caches the recent translations
    - TLB access time is part of the cache hit time
    - May allot an extra stage in the pipeline for TLB access
  - Page table storage, fault detection and updating
    - Page faults result in interrupts (precise) that are then handled by the OS
    - Hardware must support (i.e., update appropriately) Dirty and Reference bits (e.g., ~LRU) in the Page Tables
  - Disk placement
    - Bootstrap (e.g., out of disk sector 0) so the system can service a limited number of page faults before the OS is even loaded



# Check@home: Some Virtual Memory Design Parameters

|                       | Paged VM                  | TLBs                 |
|-----------------------|---------------------------|----------------------|
| Total size            | 16,000 to 250,000 pages   | 16 to 512 entries    |
| Total size (KB)       | 250,000 to 1,000,000,000  | 0.25 to 16           |
| Block size (B)        | 4,000 to 64,000           | 4 to 8               |
| Hit time              |                           | 0.5 to 1 clock cycle |
| Miss penalty (clocks) | 10,000,000 to 100,000,000 | 10 to 100            |
| Miss rates            | 0.00001% to 0.0001%       | 0.01% to 1%          |

# Check@home: Two Machines' TLB Parameters

| Characteristic   | ARM Cortex-A8   | Intel Core i7   |
|------------------|---|---|
| Virtual address  | 32 bits   | 48 bits   |
| Physical address | 32 bits   | 44 bits   |
| Page size        | Variable: 4, 16, 64 KiB, 1, 16 MiB  | Variable: 4 KiB, 2/4 MiB  |
| TLB organization | <p>1 TLB for instructions and 1 TLB for data</p> <p>Both TLBs are fully associative, with 32 entries, round robin replacement</p> <p>TLB misses handled in hardware</p> | <p>1 TLB for instructions and 1 TLB for data per core</p> <p>Both L1 TLBs are four-way set associative, LRU replacement</p> <p>L1 I-TLB has 128 entries for small pages, 7 per thread for large pages</p> <p>L1 D-TLB has 64 entries for small pages, 32 for large pages</p> <p>The L2 TLB is four-way set associative, LRU replacement</p> <p>The L2 TLB has 512 entries</p> <p>TLB misses handled in hardware</p> |

# Memory Protection

- Different tasks can share parts of their virtual address spaces
  - But need to protect against errant access
  - Requires OS assistance
- Hardware support for OS protection
  - Privileged supervisor mode (aka kernel mode)
  - Privileged instructions
  - Page tables and other state information only accessible in supervisor mode
  - System call exception (e.g., syscall in MIPS)

# The Memory Hierarchy

## The BIG Picture:

- Common principles apply at all levels of the memory hierarchy
  - Based on notions of caching
- At each level in the hierarchy
  - Block placement
  - Finding a block
  - Replacement on a miss
  - Write policy

# Block Placement

- Determined by associativity
  - Direct mapped (1-way associative)
    - One choice for placement
  - n-way set associative
    - n choices within a set
  - Fully associative
    - Any location
- Higher associativity reduces miss rate
  - Increases complexity, cost, and access time

# Finding a Block

| Associativity         | Location method                               | Tag comparisons |
|-----------------------|---|-----------------|
| Direct mapped         | Index   | 1               |
| n-way set associative | Set index, then search entries within the set | n               |
| Fully associative     | Search all entries                            | #entries        |
|                       | Full lookup table                             | 0               |

- Hardware caches
  - Reduce comparisons to reduce cost
- Virtual memory
  - Full table lookup makes full associativity feasible
  - Benefit in reduced miss rate

# Replacement

- Choice of entry to replace on a miss
  - Least recently used (LRU)
    - Complex and costly hardware for high associativity
  - Random
    - Close to LRU, easier to implement
- Virtual memory
  - LRU approximation with hardware support

# Write Policy

- Write-through
  - Update both upper and lower levels
  - Simplifies replacement, but may require write buffer
- Write-back
  - Update upper level only
  - Update lower level when block is replaced
  - Need to keep more state
- Virtual memory
  - Only write-back is feasible, given disk write latency



# Sources of Misses

- Compulsory misses (aka cold start misses)
  - First access to a block
- Capacity misses
  - Due to finite cache size
  - A replaced block is later accessed again
- Conflict misses (aka collision misses)
  - In a non-fully associative cache
  - Due to competition for entries in a set
  - Would not occur in a fully associative cache of the same total size

# Cache Design Trade-offs

| Design change          | Effect on miss rate        | Negative performance effect   |
|------------------------|----------------------------|---|
| Increase cache size    | Decrease capacity misses   | May increase access time  |
| Increase associativity | Decrease conflict misses   | May increase access time  |
| Increase block size    | Decrease compulsory misses | Increases miss penalty. For very large block size, may increase miss rate due to pollution. |

# Summary

- The Principle of Locality:
  - Program likely to access a relatively small portion of the address space at any instant of time.
    - Temporal Locality: Locality in Time
    - Spatial Locality: Locality in Space
- Page tables map virtual address to physical address
  - TLBs are important for fast translation

# Next Class

- The Processor
  - MIPS implementation
  - Pipelined processing

# Integrated Operation of the Memory System

## Computer Organization

Friday, 07 October 2022

Many slides adapted from:  
Computer Organization and Design,  
Patterson & Hennessy  
5th Edition, © 2014, MK  
and from Prof. Mary Jane Irwin, PSU



**TÉCNICO** LISBOA